

Assignment_2_HaowenShang

October 16, 2018

1 Assignment 2

1.0.1 MACS 30000, Dr. Evans

1.0.2 Haowen Shang

Due Wednesday, Oct. 17 at 11:30 AM

```
In [1]: # Import packages
import numpy as np
import pandas as pd
import statsmodels.api as sm
import matplotlib.pyplot as plt
import math
plt.style.use('seaborn')

import warnings
warnings.filterwarnings('ignore')
```

1.0.3 1. Imputing age and gender

(a) Here is where I will describe my proposed strategy ... and so on and so forth. First, we can use data in "SurveyIncome.txt" and do linear regression to find linear relationship between age, gender, total income and weight. Then, we can use these linear models and data of income and weight from "BestIncome.txt" to impute age and gender variables (total income can be calculated by labor income plus capital income). The linear models are as following:

$$\widehat{age}_i = \widehat{\alpha}_0 + \widehat{\alpha}_1 \cdot totl_inc_i + \widehat{\alpha}_2 \cdot wgt_i$$

$$P(\widehat{female}_i = 1 \mid totl_inc_i, wgt_i) = \frac{e^{\widehat{\beta}_0 + \widehat{\beta}_1 \cdot totl_inc_i + \widehat{\beta}_2 \cdot wgt_i}}{1 + e^{\widehat{\beta}_0 + \widehat{\beta}_1 \cdot totl_inc_i + \widehat{\beta}_2 \cdot wgt_i}}$$

```
In [2]: # I use the following code cells to read in my data of SurvIncome.txt, name my variable
survey_income_raw = "SurvIncome.txt"
survey = pd.read_table(survey_income_raw, sep=",", header=None)

survey_cols = [ 'total_inc', 'wgt', 'age', 'female' ]
survey.columns = survey_cols
```

```
In [3]: survey.head()
```

```
Out[3]:
```

	total_inc	wgt	age	female
0	63642.513655	134.998269	46.610021	1.0
1	49177.380692	134.392957	48.791349	1.0
2	67833.339128	126.482992	48.429894	1.0
3	62962.266217	128.038121	41.543926	1.0
4	58716.952597	126.211980	41.201245	1.0

```
In [4]: survey.tail()
```

```
Out[4]:
```

	total_inc	wgt	age	female
995	61270.538697	184.930002	46.356881	0.0
996	59039.159876	180.482304	50.986966	0.0
997	67967.188804	156.816883	40.965268	0.0
998	79726.914251	158.935050	41.190371	0.0
999	71005.223603	169.067695	48.480007	0.0

```
In [5]: survey.describe()
```

```
Out[5]:
```

	total_inc	wgt	age	female
count	1000.000000	1000.000000	1000.000000	1000.000000
mean	64871.210860	149.542181	44.839320	0.500000
std	9542.444214	22.028883	5.939185	0.500250
min	31816.281649	99.662468	25.741333	0.000000
25%	58349.862384	130.179235	41.025231	0.000000
50%	65281.271149	149.758434	44.955981	0.500000
75%	71749.038000	170.147337	48.817644	1.000000
max	92556.135462	196.503274	66.534646	1.000000

```
In [6]: # I use the following code cells to read in my data of BestIncome.txt, name my variable
```

```
best_income_raw = "BestIncome.txt"
best = pd.read_table(best_income_raw, sep=",", header=None)
```

```
best_cols = [ 'lab_inc', 'cap_inc', 'hgt', 'wgt' ]
best.columns = best_cols
```

```
best['tot_inc'] = best.lab_inc +best.cap_inc
```

```
In [7]: best.head()
```

```
Out[7]:
```

	lab_inc	cap_inc	hgt	wgt	tot_inc
0	52655.605507	9279.509829	64.568138	152.920634	61935.115336
1	70586.979225	9451.016902	65.727648	159.534414	80037.996127
2	53738.008339	8078.132315	66.268796	152.502405	61816.140654
3	55128.180903	12692.670403	62.910559	149.218189	67820.851305
4	44482.794867	9812.975746	68.678295	152.726358	54295.770612

```
In [8]: best.tail()
```

```
Out [8]:
```

	lab_inc	cap_inc	hgt	wgt	tot_inc
9995	51502.225233	14786.050723	66.781187	154.645212	66288.275956
9996	52624.117104	11048.811747	64.499036	165.868002	63672.928851
9997	50725.310645	13195.218100	64.508873	154.657639	63920.528745
9998	56392.824076	8470.592718	62.161556	145.498194	64863.416794
9999	44274.098164	12765.748454	64.974145	135.936862	57039.846618

```
In [9]: best.describe()
```

```
Out [9]:
```

	lab_inc	cap_inc	hgt	wgt	tot_inc
count	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000
mean	57052.925133	9985.798563	65.014021	150.006011	67038.723697
std	8036.544363	2010.123691	1.999692	9.973001	8294.497996
min	22917.607900	1495.191896	58.176154	114.510700	33651.691815
25%	51624.339880	8611.756679	63.652971	143.341979	61452.517672
50%	56968.709935	9969.840117	65.003557	149.947641	67042.751487
75%	62408.232277	11339.905773	66.356915	156.724586	72636.874684
max	90059.898537	19882.320069	72.802277	185.408280	98996.053756

(b) Here is where I'll use my proposed method from part (a) to impute variables.

```
In [10]: # I will use the following code cells to execute the code that will impute those vari
outcome = 'age'
features = ['total_inc', 'wgt']
X, y = survey[features], survey[outcome]
X = sm.add_constant(X, prepend=False)

m = sm.OLS(y, X)
res = m.fit()
print(res.summary())
```

```

                        OLS Regression Results
=====
Dep. Variable:          age      R-squared:          0.001
Model:                  OLS      Adj. R-squared:       -0.001
Method:                 Least Squares      F-statistic:      0.6326
Date:                   Tue, 16 Oct 2018      Prob (F-statistic):    0.531
Time:                   18:46:15      Log-Likelihood:      -3199.4
No. Observations:       1000      AIC:                6405.
Df Residuals:           997      BIC:                6419.
Df Model:                2
Covariance Type:        nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
total_inc	2.52e-05	2.26e-05	1.114	0.266	-1.92e-05	6.96e-05
wgt	-0.0067	0.010	-0.686	0.493	-0.026	0.013
const	44.2097	1.490	29.666	0.000	41.285	47.134

```
=====
Omnibus:                2.460    Durbin-Watson:                1.921
Prob(Omnibus):          0.292    Jarque-Bera (JB):        2.322
Skew:                   -0.109    Prob(JB):                0.313
Kurtosis:               3.092    Cond. No.                5.20e+05
=====
```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
 [2] The condition number is large, 5.2e+05. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [11]: ols_survey = pd.concat([y, X], axis=1)
        ols_survey.head()
```

```
Out[11]:
```

	age	total_inc	wgt	const
0	46.610021	63642.513655	134.998269	1.0
1	48.791349	49177.380692	134.392957	1.0
2	48.429894	67833.339128	126.482992	1.0
3	41.543926	62962.266217	128.038121	1.0
4	41.201245	58716.952597	126.211980	1.0

```
In [12]: ols_survey['agehat'] = res.predict(X)
        ols_survey.head()
```

```
Out[12]:
```

	age	total_inc	wgt	const	agehat
0	46.610021	63642.513655	134.998269	1.0	44.906121
1	48.791349	49177.380692	134.392957	1.0	44.545636
2	48.429894	67833.339128	126.482992	1.0	45.068980
3	41.543926	62962.266217	128.038121	1.0	44.935764
4	41.201245	58716.952597	126.211980	1.0	44.841048

```
In [13]: outcome = 'female'
        features = ['total_inc', 'wgt']
        X, y = survey[features], survey[outcome]
        X = sm.add_constant(X, prepend=False)

        m = sm.Logit(y, X)
        res = m.fit()
        print(res.summary())
```

Optimization terminated successfully.

Current function value: 0.036050

Iterations 11

Logit Regression Results

```
=====
Dep. Variable:          female    No. Observations:          1000
Model:                  Logit      Df Residuals:              997
```

```

Method:                      MLE    Df Model:                      2
Date:                        Tue, 16 Oct 2018    Pseudo R-squ.:          0.9480
Time:                        18:46:17    Log-Likelihood:         -36.050
converged:                    True    LL-Null:                -693.15
                                LLR p-value:          4.232e-286

```

```

=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
total_inc    -0.0002    4.25e-05     -3.660     0.000     -0.000    -7.22e-05
wgt          -0.4460     0.062     -7.219     0.000     -0.567     -0.325
const        76.7929    10.569      7.266     0.000     56.078     97.508
=====

```

Possibly complete quasi-separation: A fraction 0.55 of observations can be perfectly predicted. This might indicate that there is complete quasi-separation. In this case some parameters will not be identified.

```

In [14]: logit_survey = pd.concat([y, X], axis=1)
         logit_survey.head()

```

```

Out[14]:   female    total_inc      wgt  const
0        1.0  63642.513655  134.998269    1.0
1        1.0  49177.380692  134.392957    1.0
2        1.0  67833.339128  126.482992    1.0
3        1.0  62962.266217  128.038121    1.0
4        1.0  58716.952597  126.211980    1.0

```

```

In [15]: logit_survey['prob_femalehat'] = res.predict(X)
         logit_survey.head()

```

```


Out[15]:   female    total_inc      wgt  const  prob_femalehat
0        1.0  63642.513655  134.998269    1.0      0.998746
1        1.0  49177.380692  134.392957    1.0      0.999899
2        1.0  67833.339128  126.482992    1.0      0.999946
3        1.0  62962.266217  128.038121    1.0      0.999949
4        1.0  58716.952597  126.211980    1.0      0.999988

```


```

In [16]: probability_of_being_female = logit_survey['prob_femalehat']
         total_income = logit_survey['total_inc']
         logit_survey.plot(x='total_inc', y='prob_femalehat', kind='scatter')
         plt.show()

```



```
In [17]: probability_of_being_female = logit_survey['prob_femalehat']
weight = logit_survey['wgt']
logit_survey.plot(x='wgt', y='prob_femalehat', kind='scatter')
plt.show()
```



```
In [18]: def prob_female(gender):
        '''
        fuction to convert probability to dummy variable
        '''
        if gender > 0.5:
            gender = 1
        else:
            gender = 0
        return gender

logit_survey['femalehat'] = logit_survey.prob_femalehat.apply(prob_female)
logit_survey.head()
```

```
Out[18]:
```

	female	total_inc	wgt	const	prob_femalehat	femalehat
0	1.0	63642.513655	134.998269	1.0	0.998746	1
1	1.0	49177.380692	134.392957	1.0	0.999899	1
2	1.0	67833.339128	126.482992	1.0	0.999946	1
3	1.0	62962.266217	128.038121	1.0	0.999949	1
4	1.0	58716.952597	126.211980	1.0	0.999988	1

In [19]: # I will use the following code cells to execute the code that will impute those vari

```
best['imputed_age_by_inc_wgt'] = 44.2097 + 2.52e-05*best.tot_inc \
                                + (-0.0067)*best.wgt

e = math.exp(1)
best['imputed_probfemale_by_inc_wgt'] = e**(76.7929+(-0.0002)*best.tot_inc+\
(-0.4460)*best.wgt)/(1+e**(76.7929+(-0.0002)*best.tot_inc+(-0.4460)*best.wgt))
best['imputed_female_by_inc_wgt'] = best.imputed_probfemale_by_inc_wgt.\
                                apply(prob_female)

best.head()
```

```
Out[19]:
```

	lab_inc	cap_inc	hgt	wgt	tot_inc \
0	52655.605507	9279.509829	64.568138	152.920634	61935.115336
1	70586.979225	9451.016902	65.727648	159.534414	80037.996127
2	53738.008339	8078.132315	66.268796	152.502405	61816.140654
3	55128.180903	12692.670403	62.910559	149.218189	67820.851305
4	44482.794867	9812.975746	68.678295	152.726358	54295.770612

	imputed_age_by_inc_wgt	imputed_probfemale_by_inc_wgt \
0	44.745897	0.021951
1	45.157777	0.000031
2	44.745701	0.026951
3	44.919024	0.034805
4	44.554687	0.101359

```

    imputed_female_by_inc_wgt
0          0
1          0
2          0
3          0
4          0

```

In [20]: best.tail()

```

Out[20]:
      lab_inc      cap_inc      hgt      wgt      tot_inc \
9995 51502.225233 14786.050723 66.781187 154.645212 66288.275956
9996 52624.117104 11048.811747 64.499036 165.868002 63672.928851
9997 50725.310645 13195.218100 64.508873 154.657639 63920.528745
9998 56392.824076  8470.592718 62.161556 145.498194 64863.416794
9999 44274.098164 12765.748454 64.974145 135.936862 57039.846618

```

```

      imputed_age_by_inc_wgt  imputed_probfemale_by_inc_wgt \
9995          44.844042          0.004336
9996          44.702942          0.000049
9997          44.784291          0.006905
9998          44.869420          0.255027
9999          44.736327          0.991483

```

```

      imputed_female_by_inc_wgt
9995          0
9996          0
9997          0
9998          0
9999          1

```

In [21]: BestIncome = best[['lab_inc', 'cap_inc', 'hgt', 'wgt', 'imputed_age_by_inc_wgt', \
 'imputed_female_by_inc_wgt']]
BestIncome.head()

```

Out[21]:
      lab_inc      cap_inc      hgt      wgt  imputed_age_by_inc_wgt \
0  52655.605507  9279.509829 64.568138 152.920634          44.745897
1  70586.979225  9451.016902 65.727648 159.534414          45.157777
2  53738.008339  8078.132315 66.268796 152.502405          44.745701
3  55128.180903 12692.670403 62.910559 149.218189          44.919024
4  44482.794867  9812.975746 68.678295 152.726358          44.554687

```

```

      imputed_female_by_inc_wgt
0          0
1          0
2          0
3          0
4          0

```


(c) Here is where I'll report the descriptive statistics for my new imputed variables.

```
In [22]: # I will use the following code cells to do so!
        BestIncome['imputed_age_by_inc_wgt'].describe()
```

```
Out[22]: count      10000.000000
         mean         44.894036
         std          0.219066
         min          43.980016
         25%          44.747065
         50%          44.890281
         75%          45.042239
         max          45.706849
         Name: imputed_age_by_inc_wgt, dtype: float64
```

For imputed age (age_i) variable: The mean is 44.894036, the standard deviation is 0.219066, the minimum is 43.980016, the maximum is 45.706849 and number of observations is 10000.

```
In [23]: BestIncome['imputed_female_by_inc_wgt'].describe()
```

```
Out[23]: count      10000.000000
         mean         0.229400
         std          0.420468
         min          0.000000
         25%          0.000000
         50%          0.000000
         75%          0.000000
         max          1.000000
         Name: imputed_female_by_inc_wgt, dtype: float64
```

For imputed gender ($female_i$) variable: The mean is 0.229400, the standard deviation is 0.420468, the minimum is 0, the maximum is 1 and number of observations is 10000.


(d) Correlation matrix for the now six variables

```
In [24]: # Correlation matrix code and output
def corr_plot(BestIncome):
    names = BestIncome.columns
    N = len(names)

    correlations = BestIncome.corr()
    fig = plt.figure()
    ax = fig.add_subplot(111)
    cax = ax.matshow(correlations, vmin=-1, vmax=1)
    fig.colorbar(cax)
    ticks = np.arange(0,N,1)
    ax.set_xticks(ticks)
    ax.set_yticks(ticks)
    ax.set_xticklabels(names)
```

```
ax.set_yticklabels(names)
plt.show()
```

```
corr_plot(BestIncome)
```



```
In [25]: #In Matrix Form
corr = BestIncome.corr()
corr.style.background_gradient()
```

```
Out[25]: <pandas.io.formats.style.Styler at 0x1c1efd50f0>
```

1.0.4 2. Stationarity and data drift

(a) Estimate by OLS and report coefficients

```
In [26]: # Read in my third data set
income_intel_raw = "IncomeIntel.txt"
income = pd.read_table(income_intel_raw, sep=";", header=None)

# Name my variables
income_cols = ['grad_year', 'gre_qnt', 'salary_p4']
income.columns = income_cols
income.head()
```

```
Out[26]:   grad_year  gre_qnt  salary_p4
0      2001.0  739.737072  67400.475185
```

```

1      2001.0  721.811673  67600.584142
2      2001.0  736.277908  58704.880589
3      2001.0  770.498485  64707.290345
4      2001.0  735.002861  51737.324165

```

```

In [27]: # Run regression model
outcome = ['salary_p4']
features = ['gre_qnt']
X, y = income[features], income[outcome]

X = sm.add_constant(X, prepend=False)

m = sm.OLS(y, X)
res = m.fit()
print(res.summary())

```

```

                                OLS Regression Results
=====
Dep. Variable:                  salary_p4      R-squared:                0.263
Model:                            OLS      Adj. R-squared:            0.262
Method:                 Least Squares      F-statistic:                356.3
Date:                Tue, 16 Oct 2018      Prob (F-statistic):        3.43e-68
Time:                  18:46:30      Log-Likelihood:            -10673.
No. Observations:                1000      AIC:                      2.135e+04
Df Residuals:                    998      BIC:                      2.136e+04
Df Model:                        1
Covariance Type:                nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
gre_qnt	-25.7632	1.365	-18.875	0.000	-28.442	-23.085
const	8.954e+04	878.764	101.895	0.000	8.78e+04	9.13e+04

```

=====
Omnibus:                 9.118      Durbin-Watson:                1.424
Prob(Omnibus):            0.010      Jarque-Bera (JB):              9.100
Skew:                     0.230      Prob(JB):                      0.0106
Kurtosis:                 3.077      Cond. No.                      1.71e+03
=====

```


Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
[2] The condition number is large, 1.71e+03. This might indicate that there are strong multicollinearity or other numerical problems.

#Report coefficients and SE's Estimated coefficient β_0 is 8.954e+04, standard errors of β_0 is 878.764, estimated coefficient β_1 is (-25.7632), standard errors of β_1 is 1.365.

(b) Create a scatterplot of GRE score and graduation year.

```
In [28]: # Code and output of scatterplot
grad_year = income['grad_year']
gre_qnt = income['gre_qnt']
income.plot(x='grad_year', y='gre_qnt', kind='scatter')
plt.show()
```



Here is where I'll discuss any problems that jump out. I'll propose a solution here as well.

Problems: Compared data before 2011 and data after 2011, we can see that the GRE quantitative score changed a lot. In other words, the data of GRE quantitative score are drifting because scoring scale changed in 2011. Because the data of GRE quantitative score are not stationary, it can cause problems when using the ordinary least squares regression model and thus the estimated coefficients are not reliable.

Solution: We know that the GRE quantitative scoring scale is 800 before 2011 and is 170 after 2011, so we can change the data scale after 2011, which means changing the raw score in scale 170 to relative score in scale of 800 using equation ($\text{relative_score} = \text{raw_score} * 800 / 170$). In this way, the the GRE quantitative scoring scale is same before 2011 and after 2011.

```
In [29]: # Code to implement solution
income_aft2011 = income[income.grad_year >= 2011]
income_aft2011.head()
```

```
Out [29]:
```

	grad_year	gre_qnt	salary_p4
770	2011.0	148.413532	90834.606478
771	2011.0	154.123690	87255.408421

772	2011.0	155.493697	76069.366122
773	2011.0	152.551097	92160.481069
774	2011.0	156.142446	78490.139535

In [30]: `income_aft2011['gre_qnt'].describe()`

```
Out[30]: count    230.000000
mean      154.894160
std        5.197198
min       141.261398
25%       151.293028
50%       154.626456
75%       158.185442
max       170.000000
Name: gre_qnt, dtype: float64
```

In [31]: `def gre_qnt_modified(gre_qnt):`

```
    """
    function to convert GRE quantitative raw score to relative score.
    """
    gre_qnt_modified = (gre_qnt*800)/170
    return gre_qnt_modified
```

```
income_aft2011['gre_qnt_modified'] = income_aft2011.gre_qnt.apply(gre_qnt_modified)
income_aft2011.head()
```

```
Out[31]:
```

	grad_year	gre_qnt	salary_p4	gre_qnt_modified
770	2011.0	148.413532	90834.606478	698.416619
771	2011.0	154.123690	87255.408421	725.287951
772	2011.0	155.493697	76069.366122	731.735044
773	2011.0	152.551097	92160.481069	717.887515
774	2011.0	156.142446	78490.139535	734.787980

In [32]: `income_aft2011['gre_qnt_modified'].describe()`

```
Out[32]: count    230.000000
mean      728.913694
std       24.457401
min       664.759519
25%       711.967192
50%       727.653911
75%       744.402079
max       800.000000
Name: gre_qnt_modified, dtype: float64
```

In [33]: `income_bf2011 = income[income.grad_year < 2011]`
`income_bf2011.head()`

```
Out[33]:
```

	grad_year	gre_qnt	salary_p4
0	2001.0	739.737072	67400.475185

1	2001.0	721.811673	67600.584142
2	2001.0	736.277908	58704.880589
3	2001.0	770.498485	64707.290345
4	2001.0	735.002861	51737.324165

In [34]: income_bf2011['gre_qnt'].describe()

Out [34]:

count	770.000000
mean	728.421378
std	23.377864
min	655.702537
25%	712.768894
50%	727.992935
75%	744.340352
max	799.715533

Name: gre_qnt, dtype: float64

In [35]: income_bf2011['gre_qnt_modified'] = income_bf2011.gre_qnt
income_bf2011.head()

Out [35]:

	grad_year	gre_qnt	salary_p4	gre_qnt_modified
0	2001.0	739.737072	67400.475185	739.737072
1	2001.0	721.811673	67600.584142	721.811673
2	2001.0	736.277908	58704.880589	736.277908
3	2001.0	770.498485	64707.290345	770.498485
4	2001.0	735.002861	51737.324165	735.002861

In [36]: frames = [income_bf2011, income_aft2011]
income_mod = pd.concat(frames)

income_mod.head()

Out [36]:


	grad_year	gre_qnt	salary_p4	gre_qnt_modified
0	2001.0	739.737072	67400.475185	739.737072
1	2001.0	721.811673	67600.584142	721.811673
2	2001.0	736.277908	58704.880589	736.277908
3	2001.0	770.498485	64707.290345	770.498485
4	2001.0	735.002861	51737.324165	735.002861

In [37]: income_mod.tail()

Out [37]:


995	2013.0	160.441025	100430.166532	755.016586
996	2013.0	160.431891	82198.200872	754.973607
997	2013.0	154.254526	84340.214218	725.903653
998	2013.0	162.036321	87600.881985	762.523863
999	2013.0	156.946735	82854.576903	738.572869

In [38]: grad_year = income_mod['grad_year']
gre_qnt = income_mod['gre_qnt_modified']
income_mod.plot(x='grad_year', y='gre_qnt_modified', kind='scatter')
plt.show()



(c) Create a scatterplot of income and graduation year

```
In [39]: # Code and output of scatterplot
grad_year = income['grad_year']
salary = income['salary_p4']
income.plot(x='grad_year', y='salary_p4', kind='scatter')
plt.show()
```



Here is where I'll discuss any problems again ... and propose another solution.

Problems: The scatter plot above shows that the data of income 4 years after graduation are not stationary. There is an increasing trend that the mean of salary_p4 in each year increases along with the grad_year variable.

Solution: To address this problem, we need to do some linear modification of salary_p4 data. In particular, I will calculate the average salary for each graduation year, and use the following equation to standardize the salary for each year (use 2001 as the base year):

$$\text{salary_modified} = \frac{\text{salary_p4}}{(1 + \text{avg_growth_rate})^{\text{grad_year} - 2001}}$$

```
In [40]: # Code to implement a solution
#Calculate the mean salary each year
avg_inc_by_year = income['salary_p4'].groupby(income['grad_year']).\
    mean().values

#Calculate the average growth rate in salaries across all 13 years
avg_growth_rate = ((avg_inc_by_year[1:] - avg_inc_by_year[:-1]) / \
    avg_inc_by_year[:-1]).mean()

In [41]: #Divide each salary by (1 + avg_growth_rate) ** (grad_year - 2001)
income_mod['salary_modified'] = income['salary_p4'] /\
    ((1 + avg_growth_rate) ** (income['grad_year'] - 2001))
income_mod.head()
```

```
Out[41]:   grad_year   gre_qnt   salary_p4   gre_qnt_modified   salary_modified
0      2001.0   739.737072  67400.475185           739.737072      67400.475185
```


1	2001.0	721.811673	67600.584142	721.811673	67600.584142
2	2001.0	736.277908	58704.880589	736.277908	58704.880589
3	2001.0	770.498485	64707.290345	770.498485	64707.290345
4	2001.0	735.002861	51737.324165	735.002861	51737.324165

In [42]: `income_mod.tail()`

```
Out[42]:
```

	grad_year	gre_qnt	salary_p4	gre_qnt_modified	salary_modified
995	2013.0	160.441025	100430.166532	755.016586	69757.765226
996	2013.0	160.431891	82198.200872	754.973607	57094.028581
997	2013.0	154.254526	84340.214218	725.903653	58581.849117
998	2013.0	162.036321	87600.881985	762.523863	60846.675557
999	2013.0	156.946735	82854.576903	738.572869	57549.940651


In [43]: `income_mod.describe()`

```
Out[43]:
```

	grad_year	gre_qnt	salary_p4	gre_qnt_modified	\
count	1000.000000	1000.000000	1000.000000	1000.000000	
mean	2006.994000	596.510118	74173.293777	728.534611	
std	3.740582	242.361960	12173.767372	23.619014	
min	2001.000000	141.261398	43179.183141	655.702537	
25%	2004.000000	684.983551	65778.240317	712.274822	
50%	2007.000000	719.106878	73674.204810	727.910127	
75%	2010.000000	739.332537	81838.874129	744.392487	
max	2013.000000	799.715533	115367.665815	800.000000	

	salary_modified
count	1000.000000
mean	61419.808910
std	7135.610865
min	41164.726530
25%	56616.517414
50%	61467.616002
75%	66218.595876
max	84516.856633

```
In [44]: grad_year = income_mod['grad_year']
logsalary = income_mod['salary_modified']
income_mod.plot(x='grad_year', y='salary_modified', kind='scatter')
plt.show()
```



(d) Re-estimate coefficients with updated variables.

```
In [45]: # Code to re-estimate, output of new coefficients
outcome = ['salary_modified']
features = ['gre_qnt_modified']
X, y = income_mod[features], income_mod[outcome]
X = sm.add_constant(X, prepend=False)

m = sm.OLS(y, X)
res = m.fit()
print(res.summary())
```

OLS Regression Results

```
=====
Dep. Variable:          salary_modified    R-squared:                0.001
Model:                  OLS               Adj. R-squared:           -0.000
Method:                 Least Squares     F-statistic:             0.6043
Date:                  Tue, 16 Oct 2018   Prob (F-statistic):      0.437
Time:                  18:46:45          Log-Likelihood:          -10291.
No. Observations:      1000             AIC:                    2.059e+04
Df Residuals:          998              BIC:                    2.060e+04
Df Model:              1
Covariance Type:       nonrobust
=====
```

coef	std err	t	P> t	[0.025	0.975]
------	---------	---	------	--------	--------

```

-----
gre_qnt_modified    -7.4321      9.560      -0.777      0.437      -26.193      11.329
const              6.683e+04   6968.684      9.591      0.000      5.32e+04   8.05e+04
=====
Omnibus:              0.789   Durbin-Watson:              2.025
Prob(Omnibus):        0.674   Jarque-Bera (JB):              0.698
Skew:                 0.060   Prob(JB):              0.705
Kurtosis:             3.050   Cond. No.              2.25e+04
=====

```

Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 2.25e+04. This might indicate that there are strong multicollinearity or other numerical problems.

Here is where I'll discuss how the coefficients differ, where I'll interpret why the changes result in new coefficient changes, and where I'll discuss what this suggests about the answer to the question.

Estimated coefficient β_0 is 6.683e+04, standard errors of β_0 is 6968.684, estimated coefficient β_1 is (-7.4321), standard errors of β_1 is 9.560.

The estimated coefficient β_0 is still near zero, and the standard error of β_0 increases by 8 times. The absolute value of coefficient β_1 is just 30% of the previous one because we changed the scale of part of gre_qnt data (increase the scale of data after 2011) and we eliminated potential time trend in the salary_p4 data. The standard errors of β_1 increased by 7 times. Compared to previous OLS results, we find that β_1 is less negative and isn't statistically significant, which is much more reliable than previous significant result.

The p-value is large and the coefficient β_1 is not statistically significant, so the OLS regression result gives no evidence that higher intelligence is associated with higher income.

1.0.5 3. Assessment of Kossinets and Watts.

Please see attached PDF.

Assessment of Kossinets and Watts (2009)

The fundamental question the paper is trying to answer is “What are the mechanisms of homophily emergent over time based on the decisions of individuals to make and break ties?”

In order to answer this question, the author did the analysis based on the population of 30,396 undergraduate and graduate students, faculty, and staff in a large U.S. university, who used their university e-mail accounts to both send and receive messages during one academic year. The data set comprised 7,156,162 messages exchanged by 30,396 stable e-mail users during 270 days of observation. The sources of the data set are **three different databases**: (1) the logs of e-mail interactions within the university over one academic year, (2) a database of individual attributes (status, gender, age, department, number of years in the community, etc.), and (3) records of course registration, in which courses were recorded separately for each semester. In the analysis, the number of observations are **30396** for personal characteristics, organizational affiliations, and course-related variables; the number of observations are **7,156,162** for email messages; and the time period of the data is **270 days**. The author listed the description and definition of all the variables in '**APPENDIX A**'.

However, I find a potential problem in the data cleaning process, which is that only e-mail accounts on the central university server were included in the data set. But in fact, a number of individuals also used accounts provided by their departments, such as xyz@department.university.edu (mostly the faculty and graduate students in departments such as computer science, mathematics, and physics). The departments of computer science, mathematics, and physics provide basic courses which most students would be involved in, so people who use e-mail accounts provided by departments have extensive interaction with other people. Excluding these data will lose some information and sources of homophily.

Besides the data cleaning process, there is also one weakness of matching theory variable to data variable. In particular, in this paper, theory variable is “social relationships” and data variable is e-mail logs linked to other characteristics of the senders and receivers. In order to understand how homophily emerges over time as a function of the decisions of individuals to

make and break ties, the author focus highly on the formation of new ties. One kind of tie formation mechanisms is 'focal closure', which is defined as the various groups, contexts, and activities around which social life is organized and which in turn facilitate interpersonal interactions. When matching the 'focal closure' to the data set, the author is supposed to have a record of all possible focal activities, including classes, social groups, sporting and cultural organizations, shared housing, and so forth—so that he could study separately their effects on social interactions over time. However, the data set can just provide classes administered by the university registrar as explicit foci, which are certainly not the only foci of interaction. The author overcome this practical obstacle in part by mining the available data in more creative ways. Specifically, he makes use of the “bulk” messages which are defined as having more than one recipient and are used to infer the presence of shared groups and activities that were not otherwise recorded in the data set. The author treats bulk messages as indicators of social foci, defined broadly as any kind of shared affiliation, group, or activity that generates a demand for group-oriented communication. In this way, there is a much better match between focal closure and the dataset.